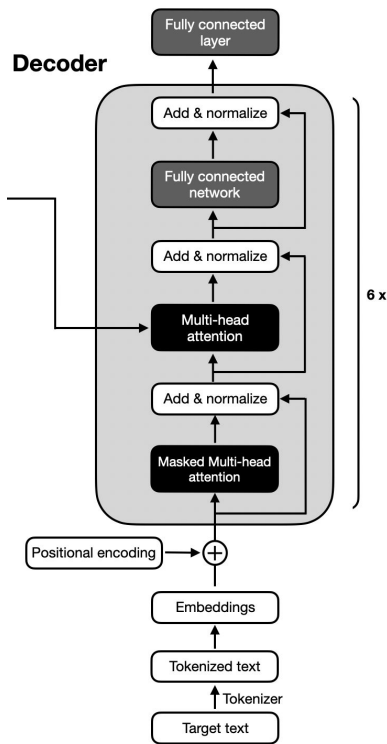

LLM interpretability via sparse autoencoders

INFO 3350/6350: Lecture 20

Why does interpretability matter?

- Can you trust your model?
 - Can you trust your data?
 - Is the model biased or otherwise harmful?
 - Are the features important?
 - Are the features useful?
 - Training and evaluation costs are high
 - Deployment stakes can be high
-



Approaches to LLM interpretability

Weights are hopeless, but *neurons* might be useful.

There are a lot of neurons, but far fewer than there are total params. Each neuron has many weights and neurons are organized into layers.

Give input, see which neurons activate at a given layer. Hope that there is one (“monosemantic”) neuron that fires. Maybe look for “circuits” of neurons that fire across layers.

Works, sometimes, but polysemanticity and superposition ...

Superposition

Superposition

Ideally, want an interpretable neuron to activate iff a single “feature” is present in the input

Aside: what’s a feature? Coherent concept *in the data*, nothing to do with the model

Some neurons activate in the presence of multiple, *disparate* features (e.g., cats, hex numbers, gerunds). And some features activate many neurons, but only weakly. But not always ...


[Elhage et al., “Toy Models of Superposition” \(2022\)](#)

Superposition


Superposition means that features are represented via combinations of neuron activations. The activating neurons *do not* necessarily have anything to do with the feature in question.

Superposition is *required* to represent more features than there are neurons in the model (e.g., all people, ideas, etc.)

Some features are more important than others



Polysemanticity is what we'd expect to observe if features were not aligned with a neuron, despite incentives to align with the privileged basis.



In the **superposition hypothesis**, features can't align with the basis because the model embeds more features than there are neurons. Polysemanticity is inevitable if this happens.

Example of superposition

In a small language model ([“Toward monosemanticity”](#)), there are features for Arabic script and for Hebrew-language text. But no single neuron activates strongly on all Arabic or Hebrew text input. In fact, the top-activating examples for the neurons that *do* fire on such input contain no examples of Arabic or Hebrew inputs.

Clearly, these features are represented only via superposition. They exist across low levels of activation for many neurons.


Superposition

Superposition only works well if features are *sparse*, i.e., if most inputs have few features, and if the network has a nonlinear activation function.

If these conditions are not met, superposition causes destructive interference, because you can't tell whether the multiple activations represent a single, superposed feature or a mix of multiple features.

Linear Model


(or any)



Linear models learn the top m features. $1 - S = 0.001$ is shown, but others are similar.


ReLU Output Model

$1 - S = 1.0$




In the **dense** regime, ReLU output models also learn the top m features.

$1 - S = 0.3$





As **sparsity increases**, superposition allows models to represent more features. The most important features are initially untouched. This early superposition is organized in antipodal pairs (more on this later).


$1 - S = 0.1$




$1 - S = 0.03$




$1 - S = 0.01$



$1 - S = 0.003$



$1 - S = 0.001$



Weight / Bias
Element
Values
-1 0 1

Superposition
 $\sum_j (\hat{x}_i \cdot x_j)^2$
0 1

Parameters
 $n = 20$
 $m = 5$
 $I_i = 0.7^i$

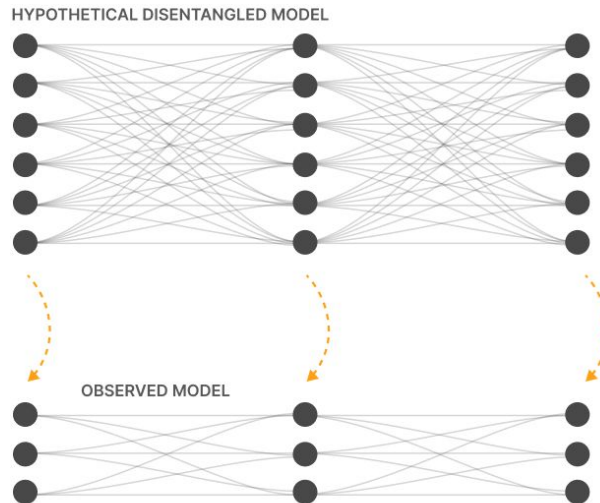


Note
this

Hypothetical networks

Sparse networks with superposition can allow a network to function as a simulation of a larger network that is fully *disentangled*

LLMs may be such networks *and* we can use the same hypothesis to *probe* LLMs



Under the superposition hypothesis, the neural networks we observe are **simulations of larger networks** where every neuron is a disentangled feature.

These idealized neurons are **projected** on to the actual network as “almost orthogonal” vectors over the neurons.

The network we observe is a **low-dimensional projection** of the larger network. From the perspective of individual neurons, this presents as polysemanticity.


Sparse autoencoders

SAE architecture

Begin with values from hidden state of a layer of neurons for many different inputs

Seek to learn a *sparse* representation in *higher-dimensional* space that reconstructs the inputs with minimal loss

Need to penalize non-sparsity, or we just get pseudo-identity



SAE interpretation

Hidden layer in the SAE is the hypothesized large feature set.
In *this space*, 1 neuron = 1 feature.

For a given input, the feature set is manageable for interp, because most features/activations are zero.

Can find inputs that maximally activate each neuron/feature in the SAE.

Can change values of SAE neurons and observe effect on output logits (i.e., next token predictions).

Feature Number
(click for hyperlink)

Human
explanation

Histogram of randomly
sampled non-zero
activations

Top 10 negative and
positive output logits of
the feature

Top 20 max
activating examples

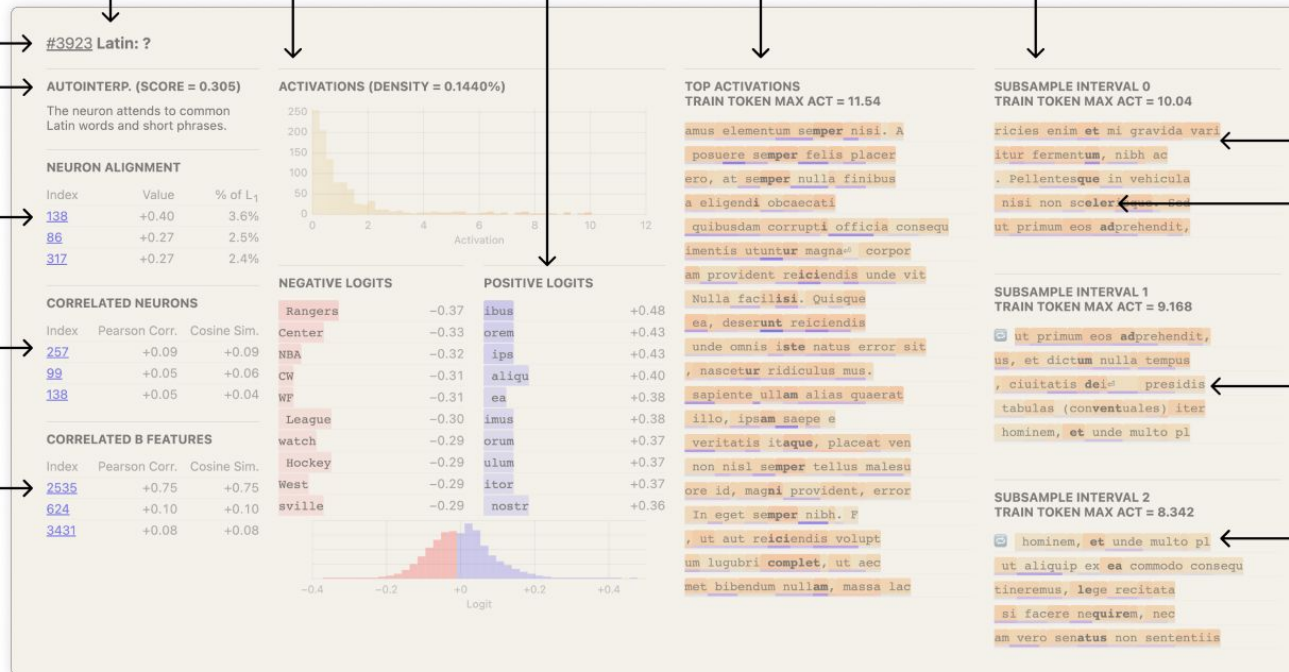
Ten evenly spaced intervals
spanning the full range of
activation values

Autointerp
explanation and
prediction score

Top 3 neurons
by how much the
feature activates them

Top 3 neurons
by token correlation


Top 3 features from
the parallel run with
a different random
seed



Blue underline means a
lower ablation loss
(better token prediction);
red means a higher loss

Bold token comes from
the training data (used to
select each example).
Surrounding 4 tokens
give context

Hover over any token
(for 2+ seconds) to see
its activation value and
ablations

 means this specific
example has already
appeared in another
interval

Evaluating SAE features

Are the features that we find meaningful and monosemantic?

Are similar features found across multiple models and SAE settings?

Do features split as the size of the hidden layer in the SAE increases and coalesce as it decreases?

Can SAE features be intervened on in predictable ways?

Feature Number
(click for hyperlink)

Human
explanation

Histogram of randomly
sampled non-zero
activations

Top 10 negative and
positive output logits of
the feature

Top 20 max
activating examples

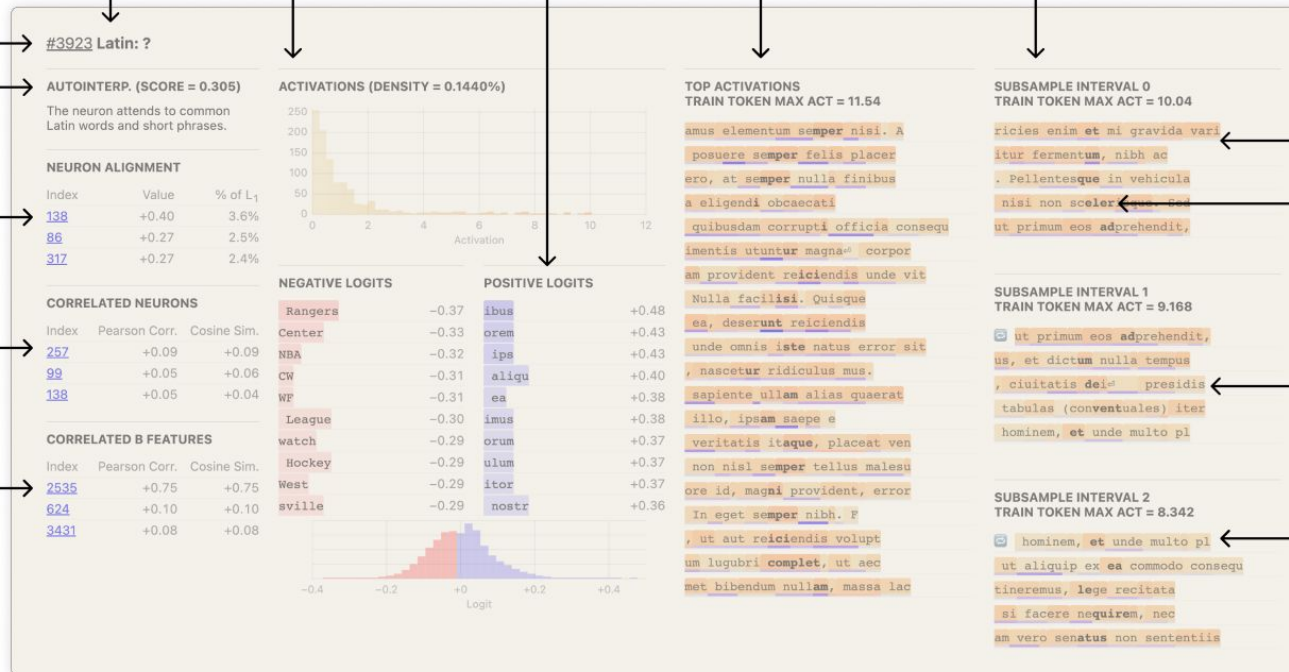
Ten evenly spaced intervals
spanning the full range of
activation values

Autointerp
explanation and
prediction score

Top 3 neurons
by how much the
feature activates them

Top 3 neurons
by token correlation


Top 3 features from
the parallel run with
a different random
seed



Blue underline means a
lower ablation loss
(better token prediction);
red means a higher loss

Bold token comes from
the training data (used to
select each example).
Surrounding 4 tokens
give context

Hover over any token
(for 2+ seconds) to see
its activation value and
ablations

 means this specific
example has already
appeared in another
interval